

Detection of human alertness using supervised learning

Anders Hørsted

Kongens Lyngby 2011
IMM

Technical University of Denmark
Informatics and Mathematical Modelling
Building 321, DK-2800 Kongens Lyngby, Denmark
Phone +45 45253351, Fax +45 45882673
reception@imm.dtu.dk
www.imm.dtu.dk

Summary

This is the summary/abstract

Resumé

På dansk

Preface

This thesis was prepared at Informatics Mathematical Modelling, the Technical University of Denmark in partial fulfillment of the requirements for acquiring the Ph.D. degree in engineering.

The thesis deals with different aspects of mathematical modeling of systems using data and partial knowledge about the structure of the systems. The main focus is on extensions of non-parametric methods, but also stochastic differential equations and neural networks are considered.

The thesis consists of a summary report and a collection of ten research papers written during the period 1996–1999, and elsewhere published.

Lyngby, December 1999

Henrik Aalborg Nielsen

Acknowledgements

I thank my...

Contents

I	Introduction and data exploration	1
1	Introduction	3
1.1	The competition	4
1.1.1	Problemformulering	4
2	The Ford Challenge	5
2.1	Other online machine learning competitions	5
2.1.1	Netflix Prize	6
2.1.2	KDD Cup	6
2.1.3	And many others	7
2.2	Kaggle.com and The Ford Challenge	7
2.2.1	The data set	9
3	Data exploration	11
3.1	A note about Test data, training data	11
3.2	Calculating common statistics	11
3.3	Determining the datatype of features	12
3.3.1	Unique values	13
3.3.2	Plotting some features	15
3.4	Outlier detection	17
3.4.1	Making boxplots of features	17
3.4.2	Making layered feature plots	19
3.4.3	Features that are constant in a trial	20
3.5	Finding possible discriminating features	21
3.5.1	Testing binary features	21

3.5.2	Scatterplots	21
3.6	Principal Component Analysis	22
3.6.1	Theory	22
3.7	Conclusions	23
II	Modelling	25
4	Theory of classification	27
4.1	The binary classification problem	27
4.2	Parametric models and maximum likelihood	29
4.2.1	Maximum likelihood	29
4.3	Logistic Regression	30
4.3.1	The decision boundary	31
4.3.2	Model complexity	32
4.4	Feed Forward Neural Network	33
4.4.1	Model complexity	33
4.4.2	Regularization	33
4.5	Measuring classifier performance	33
4.5.1	Splitting the dataset	34
4.5.2	Cross validation	35
4.5.3	AUC	36
4.6	Model selection	36
4.6.1	Over- and underfitting	36
4.6.2	Forward feature selection	36
4.6.3	Regularization	36
5	Recreating winning approach	37
5.1	The winning approach	37
5.1.1	Classification method	38
5.2	Recreating the winning approach	40
5.2.1	Statistics on the AUC score	41
6	Improving the winning approach	43
6.1	Forward selection	43
6.1.1	Testing performance on the original testset	45
7	Other classification methods	47

III	Workflow and discussion	49
8	Workflow and tools	51
9	Discussion	53
10	Conclusion	55
A	Appendices	57
A.1	Calculating common statistics - Code	58
A.2	Calculating common statistics - Result	59
A.3	Calculate unique values of features - Code	60
A.4	Calculate unique values of features - Result	62
A.5	Creating boxplots - Code	62
A.6	Creating boxplots - Result	63
A.7	Creating layered featuer plots - Code	63
A.8	Creating layered feature plots - Result	63
A.9	Scatterplots - Code	63
A.10	Scatterplots - Result	64
A.11	Recreating the winning approach - Code	64
A.12	Forward selection - Code	65
A.13	Forward selection - Results	67

Part I

Introduction and data exploration

CHAPTER 1

Introduction

Igennem de sidste 30 år er computerens ydeevne vokset markant. Den forbedrede ydeevne har givet mulighed for at udføre statistisk dataanalyse, i et omfang der ikke tidligere har været muligt. En af de ting der er blevet mulighed for, er at programmere såkaldte "classifiers". Et godt eksempel på en classifier, er de programmer bankerne bruger til at opdage evt. snyd med kreditkort. Baseret på alle tidligere korttransaktioner, og viden om hvilke der var "falske" transaktioner, kan bankerne programmere en classifier, der med høj præcision kan forudsige om en ny transaktion er snyd eller ej.

Denne opgave tager udgangspunkt i en konkurrence på hjemmesiden [kaggle.com](https://www.kaggle.com). Konkurrencen er udbudt af Ford Motors og handler om at udvikle en classifier, der kan forudsige om en bilist er ved at blive ukoncentreret, mens han/hun kører bil. Til at udvikle denne classifier har Ford foretaget en række målinger på bilister, mens de kørte bil, og til hver måling er det blevet noteret om bilisten var opmærksom eller ej. Med udgangspunkt i disse målinger udarbejdes – ved brug af de mest gængse metoder – en række classifiers, og deres evne til at klassificere ny data sammenlignes.

1.1 The competition

1.1.1 Problemformulering

I denne opgave vil jeg...

- ... bruge de mest gængse klassifikationsmodeller (nearest neighbour, logistic regression, neural networks og SVM) til at lave en classifier der (forhåbentlig) kan forudsige om en bilist er ved at falde i søvn.
- ... undersøge hvor stor indflydelse den indledende databehandling (feature selection og outlier removal) har på det endelige resultat.
- ... undersøge om classifieren kan forbedres ved at implementere en Hidden Markov Model, der tager hensyn til det temporale aspekt af data.
- ... implementere en ensemble classifier, der kombinerer resultatet af flere classifiers i én classifier.

CHAPTER 2

The Ford Challenge

In this chapter, the ford competition that is the basis for this report, is introduced. Before introducing the ford challenge, a short overview of other online machine learning competitions is given. As part of introducing the ford competition, the kaggle.com website that hosted the ford competition is presented, and the data set used in the ford competition is described in detail. But as mentioned the chapter starts with a short overview of other online machine learning competitions.

2.1 Other online machine learning competitions

To get a little perspective on the ford challenge, this section gives a short review of some past and present online machine learning competitions.

2.1.1 Netflix Prize

One of the most talked about competitions, may very well be the Netflix Prize. The Netflix competition was launched on October 2, 2006 and the aim of the competition was to predict how users would grade new movies, based on a large dataset of previous grades. Why did the Netflix Prize gather a lot of attention? First of all the grand prize was \$1M, which is a lot of money. And secondly the dataset was huge, consisting of 100,480,507 ratings given by 480,189 users, leaving room for a lot of interesting new techniques to be tested.

When the Netflix Prize was awarded on September 18, 2009, 5169 different teams had submitted at least one entry to the competition. The winning team consisted of three different teams, that at one point decided to team-up and compete as a join-team.

Netflix originally wished to follow up the Netflix Prize, with another competition but decided to dismiss the idea, due to a lawsuit regarding privacy concerns related to the first Netflix Prize. *

2.1.2 KDD Cup

Although the Netflix Prize gather a lot attention, it wasn't the first online machine learning competition. An example of an earlier competition, is the KDD Cup. The KDD Cup is a competition that is held every year, and that started back in 1997. The subject changes every year, and can be anything from mining purchase data from an online store, to computer aided detection of breast cancer (the 2000 and 2008 competitions respectively).

This year the KDD Cup is held in cooperation with Yahoo! Labs and the task is to predict user ratings of musical items (both tracks, albums, artists and genres). One note worthy detail about the 2011 KDD Cup is the huge data set, containing over 300 million ratings of more than 600,000 distinct items. †

* The section about The Netflix Prize is based on [Wikipedia \(2011\)](#) and [Netflix \(2011\)](#)

† Read more about the KDD Cup history at [ACM \(2011\)](#). For more info about the 2011

2.1.3 And many others

The Netflix Prize and the KDD Cup are just two examples of online machine learning competitions. Many others exist, such as

- *The Hearst Challenge 2011* - Every year The Hearst Corporation hosts a machine learning competition. This year the task is to data mine the history of 1.8 million emails sent to subscribers of Hearst's publications, and then predict who will open emails in the future. Read more at <http://www.hearstchallenge.com>
- *The Reclab Prize* - RichRelevance is a company that specializes in online product recommendation. They offer a \$1M prize for the team that first improves their product recommendation algorithm by 10%. Read more at <http://www.overstockreclabprize.com>
- *The Heritage Health Prize* -

Since almost all the machine learning competition websites, need the same functionality, websites that specialize in hosting machine learning competitions have appeared. One example of a website that provides a hosting platform for machine learning competitions is the kaggle.com website, who hosts The Ford Challenge.

2.2 Kaggle.com and The Ford Challenge

TODO: Mention that use of physical features should be minimized?

As already mentioned, kaggle.com hosts machine learning competitions for universities and corporations. The first competition hosted by kaggle.com was started in April 2010, and since then a total of 18 competitions have been held. Every competition at kaggle.com has some background information, links to the data sets, a submission system and a forum, as seen in figure 2.1.

competition see [Labs \(2011\)](#)




Figure 2.1: The Ford Challenge frontpage at the kaggle.com website

The Ford Challenge began on January 19, 2011. The task was to create a classifier that is able to detect when a driver is about to get distracted while driving. The dataset that Ford made available for the competition consisted of measurements of 30 different features, measured on drivers along with a binary feature (IsAlert) that was 1 if the driver was alert and 0 otherwise. The 30 features was a mix of environmental, driver physiological and vehicular features[‡]. Based on this dataset a classifier should predict the IsAlert feature of a distinct test dataset held by Ford.

One detail that made the competition slightly different than many other competitions, was that Ford would not disclose any information about what the different features represented[§]. The official reason was that (Abou-Nasr, 2011c)

“We like to encourage the participants to pursue classification without preconceived notions based on prior knowledge of the

[‡]The dataset is explained in much more detail in the section 2.2.1

[§]see forum replies from the Ford spokesperson (Abou-Nasr, 2011b,a)

subject, focusing on variables which lead them (based on their own experiments) to better classification."

Doubts about the true motive behind the lack of details about the features, was expressed by, what later turned out to be, the winner of the competition ([Inference, 2011b](#))

The performance of the classifiers was measured by calculating the AUC (see section 4.5.3) of the classifiers, on the test set. A limit of two submissions per contestant was set as a way to counteract the possibility of someone reverse-engineering the IsAlert-feature of the test dataset. This could of cause be circumvented by one person registering more than once. And this was exactly what happened. On March 9, 2011 when the competition ended, two users had achieved exactly the same AUC (six significant digits) and other contestants immediately questioned the probability of two unrelated users getting exactly the same AUC ([Pardos, 2011](#)). One of the two leaders (Rosanne/Shen) quickly admitted that he and his friend had indeed used two accounts to get 4 submission per day ([Rosanne, 2011](#)), and after some discussion in the forum, the leaders were disqualified. The end result was that the user Inference in third position was declared the official winner of the competition.

Two weeks after the competition ended, Inference described the technique used to win the competition. This will be described in detail in chapter 5. But before that chapter, the data set will be described in details in the next section.

2.2.1 The data set

The dataset used to create a classifier was released on the competition website, the day the competition started. The dataset consisted of a number of trials and each trial was approximately 2 minutes of sequential data recorded every 100ms during a driving session on the road or in a driving simulator ([Kaggle.com, 2011](#)). As the interval between two rows was 100ms, each trial consists of approximately $2\text{minutes} \cdot 60 \frac{\text{secs}}{\text{minute}} \cdot 10 \frac{\text{rows}}{\text{sec}} = 1200$ rows.

TrialID	ObsNum	IsAlert	P1	...	P8	E1	...	E11	V1	...	V11
0	0	0	12.2	...	1.2	4.3	...	33	12	...	7.34
⋮											⋮
0	1200	1	11.1	...	10.7	1.3	...	21	8	...	8.82
⋮											⋮
510	1198	0	11.1	...	10.7	1.3	...	21	8	...	8.82

Table 2.1: Structure of the data set

Each row has a total of 33 data columns structured as shown in table 2.1. Some details are worth noticing:

- The TrialID starts at 0 and the trials from 469 to 479 (both inclusive) are missing. The last trial has TrialID=510. This gives a total of exactly 500 trials.
- The ObsNum also starts at 0 there are not exactly 1200 observation for every trial.
- The total number of rows is 604,229
- The row number is not part of the data set, so a row is uniquely identified by the pair (TrialID, ObsNum).

As mentioned before no additional information about what the different features represent or what datatype (discrete, continous) they are, was disclosed by Ford. The only way to get these informations is by doing a thorough data exploration of the data set and that is what the next chapter is about.

CHAPTER 3

Data exploration

Here I describe the various data exploration techniques I have used. Lots of nice graphs.

3.1 A note about Test data, training data

3.2 Calculating common statistics

To start the data exploration, four common statistics, namely the mean, min, max and standard deviation, of every feature across the whole dataset was calculated. The standard deviation was calculated as (with n equal the total number of rows in the dataset, and x_i equal to the i 'th value of any of the features)

$$s = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n}}$$

Feature	Mean	Min	Max	Std.Dev
E9	0.87	0.00	1.00	0.33
IsAlert	0.58	0.00	1.00	0.49
P6	843.73	128.00	228812.00	2795.32
P8	0.00	0.00	0.00	0.00
V5	0.18	0.00	1.00	0.38
V7	0.00	0.00	0.00	0.00
V9	0.00	0.00	0.00	0.00

Table 3.1: Highlights from the results of the summary statistics. See appendix A.2 for all results.

The source code for the calculations can be found in appendix A.1 and all results in appendix A.2. Most results did not tell that much, but a few results stood out. These are shown in table 3.1. The table shows that the features P8, V7 and V9, are zero throughout the whole dataset, and can be ignored. The features E9 and V5, could be binary, but that is only speculation at the moment. Finally the feature P6 has a mean of 843.73 and a standard deviation of 2795.32, but its maximum is 228812.00, which is many, many standard deviations away from the mean. This could be a sign of some serious outliers, but it could as well be a single trial with a mean far from the other trials. Further investigation are needed to conclude anything here. Finally it is seen that the mean of the IsAlert feature is only a little above 0.5, and therefore only a little over 50% of the time are the drivers alert. This may be a bit surprising.

The results of this first step have been to exclude a few features, and get some rough ideas about the shape of some other features. It is now time to resally get to know the different features.

3.3 Determining the datatype of features

Since Ford would not disclose any information about the different features, it is important to get a good picture of which features are discrete/-categorical and which features are continuous. A natural first step to learn the datatype of the features, is to calculate the number of unique values each feature takes.

3.3.1 Unique values

It requires a little thought to pinpoint exactly what needs to be calculated. On one hand it is natural to calculate the number of unique values a feature takes across the whole dataset. On the other hand it could happen, that a discrete feature might have only (eg.) 3 unique value within any trial, but the values differ between various trials. This way we would have a feature with 1500 unique values across the whole dataset, but with max 3 unique values within any given trial. Therefore the number of unique values within a trial are calculated for every feature and every trial. Based on this result, the minimum and maximum number of unique values within a single trial, is calculated for every feature. The source code can be found in appendix A.3 and results are shown in table 3.2. **TODO: The result table in appendix or in report?**

A couple of things should be noticed about the results. Almost all features have som trials where they only take on one unique value. For categorical features, this could be perfectly normal, but for a continous feature it seems to be pretty unlikely to have only one value in a trial spanning two minutes. Is a feature, that in some trials seems continous, but then only have one value in other trials, just turned off in the latter trials? And how should the feature be handled in trials where it is “off”? This discussion is continued in section 3.4 about outlier detection.

From table 3.2 it is also seen, that only features P1, P2 and V11 is consistently having lots of unique value, across all trials. It seems fair to call these features continous. Feature V1 have some trials where it takes on 969 different values, but it also have trials with only one unique value. Across the whole dataset it takes on 12374 unique values. This could be explained by feature V1 being continous in a small number of trials (about 15-30), and turned “off” in all other trials. Further exploration will show whether it is true or not.

Another interesting detail is that there are two pairs of features ((P3,P4), (P6,P7)) in the results, that share exactly the same number of unique features. Both min and max and total. This indicates a possible relationship between the features, and later (section 3.5.2) it is shown that this is in fact true.

Feature	Min in trial	Max in trial	Whole dataset
E1	1	131	12265
E10	1	86	121
E11	1	79	116
E2	1	129	28502
E3	1	3	3
E4	1	97	249
E5	1	134	254
E6	1	116	247
E7	1	15	25
E8	1	10	10
E9	1	2	2
IsAlert	1	2	2
P1	967	1208	105715
P2	1155	1207	128666
P3	24	106	406
P4	24	106	406
P5	34	114	1070
P6	3	99	406
P7	3	99	406
V1	1	969	12374
V10	1	5	5
V11	875	1206	166455
V2	1	76	90
V3	1	29	33
V4	1	204	324
V5	1	2	2
V6	1	775	2727
V8	1	239	323

Table 3.2: The minimum and maximum number of unique values within the trials, for every feature in the dataset. Also the total number of unique values for each feature, across the whole dataset, are shown.

A final remark about the number of unique values, is that feature E9 and V5 indeed are binary. Also if a limit of max 40 unique values within a single trial is set, as an indicator for categorical features, E3, E7, E8, V3 and V10 are seen to be categorical.

3.3.2 Plotting some features

To get a more nuanced picture of the various features, it is now time to visualize the data. The first thing of interest is to plot every feature for a couple of trials, to see the structure of the features, and how much they vary between different trials. For a start some random plotting was done in the interactive ipython-shell*.

Some selected plots from the interactive plotting session are shown in figure 3.1. The main results from the introductory plotting of features are that:

- A feature varies a lot in structure between different trials
- Many features are constant in some trials (was already hinted at in section 3.3.1)
- Some features deviate a lot from their “normal” structure, in a few trials (more about that in section 3.4 about outliers).
- The feature P2 looks like pure noise **TODO: Is this right?**
- The features E3, E7, E8, V3 and V10 do indeed seem to be categorical features, although V3 does have some extra structure **TODO: What do I mean? Plot categorical features?**

Apart from the observations above, a subjective conclusion from the plots of the features is that the data quality could be better. Since the data quality isn’t perfect it is important to consider how outliers should be dealt with. This is what the next section is about.

* For details about the software setup used, see chapter 8

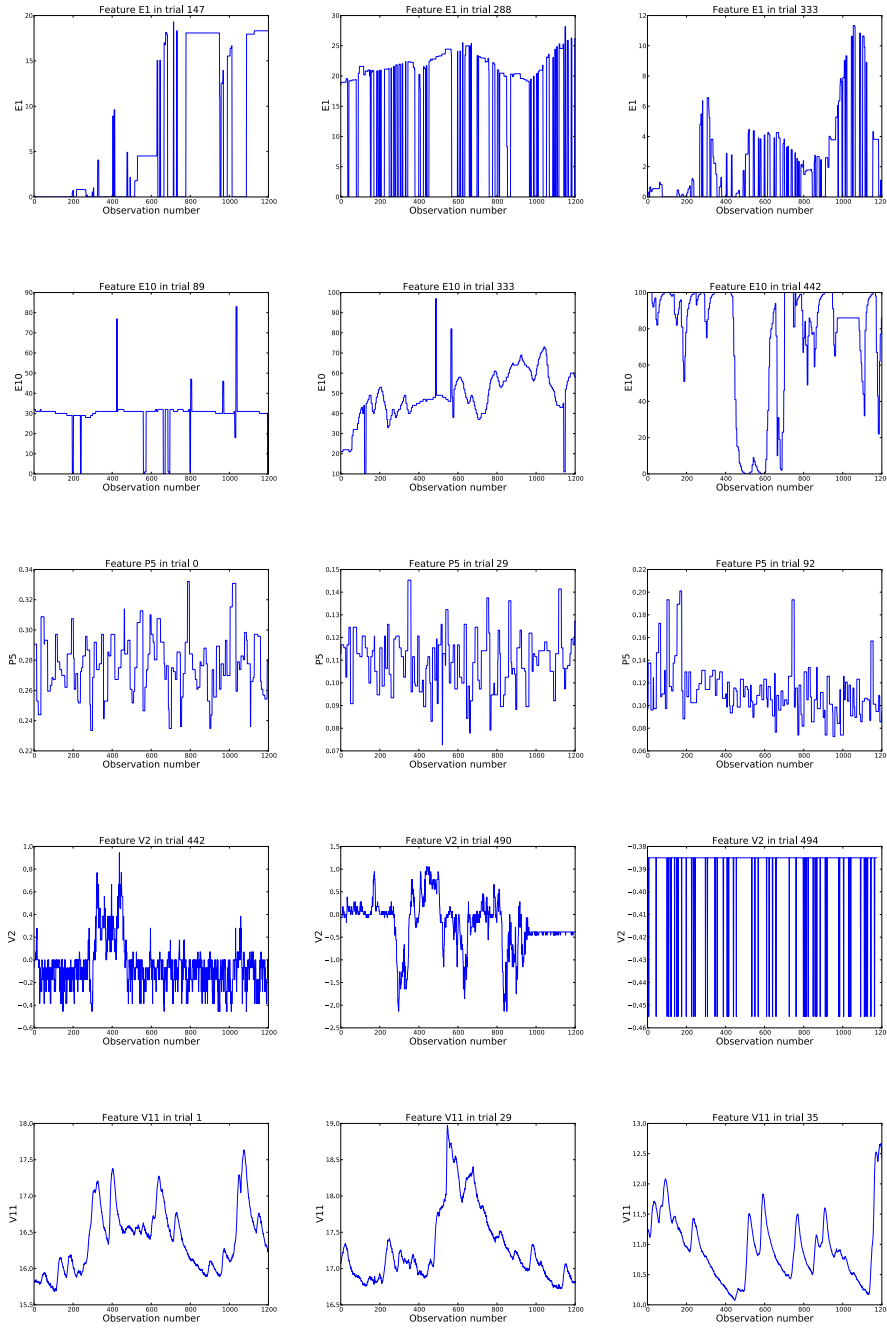


Figure 3.1: Plots of the features E1, E10, P5, V2, V11, for a couple of trials. The trials were selected to show the diversity a single feature exhibits between trials.

3.4 Outlier detection

The textbook approach for detecting outliers is to create boxplots of the different features, and then argue that data points outside the 95-percentile should be removed. Although this is a naive approach some useful information could be obtained from boxplots of the features. But it doesn't necessarily make sense to make boxplots of the features for the whole dataset. Some features have the same structure in different trials, but their mean varies. Eg. feature V11 in trials 29 and 35 have similar structure but the range in trial 29 is [16.75, 19] and [10, 12.6] in trial 35 (see figure 3.1). If a feature has a different mean in only a few trials, those datapoints could be seen as outliers, if we look at the feature across the whole dataset. Instead boxplots of the features for the 50 first trials are created. This also gives a visual idea about how feature means varies from trial to trial.

3.4.1 Making boxplots of features

The boxplots for 4 selected features are shown in figure 3.2. All boxplots are in appendix A.6, and the source code can be seen in appendix A.5.

From the boxplots a couple of things are of interest.

- For some features the trial mean varies a lot between different trials
TODO: Remember to check with boxplots in appendix
- Within a single trial some features still have a lot of data points outside the 95 percentile. See eg. P1 in figure 3.2 **TODO: Check the right name (percentile?)**

As many features have datapoints outside the 95 percentile even within a single trial, the next step could be to categorize these datapoints as outliers. But inspection of some examples of specific features in specific trials, shows that it would be wrong to remove these "outliers". Eg. feature V11 in trial 29 is plotted in figure 3.1 and seems to be perfectly normal. Careful inspection of the boxplot for feature V11 in trial 29 (figure 3.2) shows that the feature has many outliers in this specific trial.

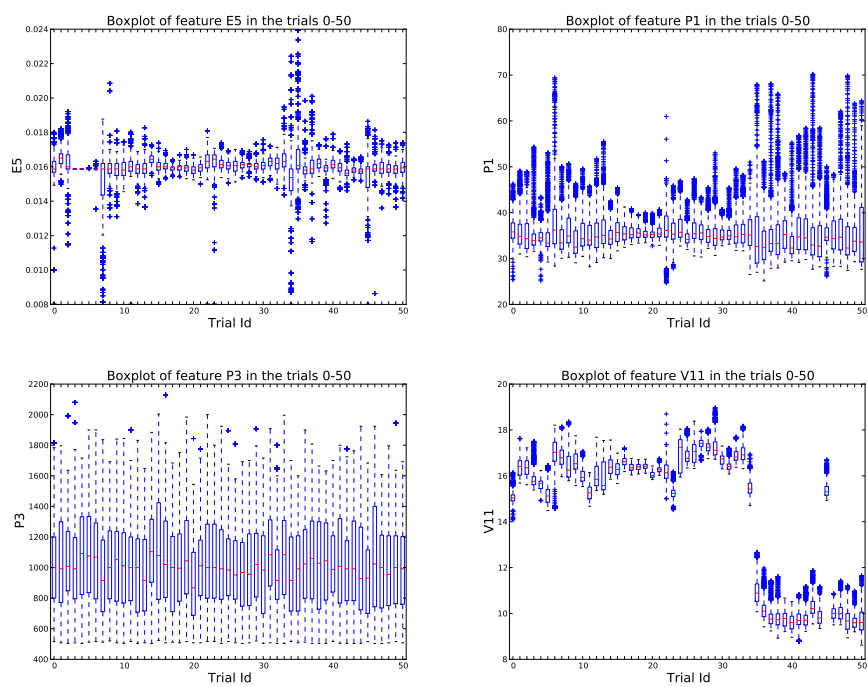


Figure 3.2: Boxplots of feature E5, P1, P3 and V11 in the first 50 trials




Figure 3.3: Layered feature plots of feature P6. All 500 trials are shown on the left. A couple of possible outliers are seen. On the right only 20 randomly chosen trials are plotted

The boxplots do not correctly identify the outliers in the dataset, and another method is therefore needed.

3.4.2 Making layered feature plots

Instead of trying to categorize single datapoints as outliers, it may make more sense to define a specific feature in a specific trial as an outlier. The fact that some features deviates a lot from their normal structure, in a few trials was already mentioned in section 3.3.2 about exploratory plotting of features. Since the exploratory plotting was a rather unstructured process, a more thorough picture, of the structure of the different features across all trials, is needed. One way to get a better picture is by plotting all trials of a single feature, in a single plot. The graph of each trial is given a low opacity and then all 500 trials **TODO: update number based on number of training trials** are plotted on top of each other. That way the common structure of a feature across trials are visualized, and possible outlier trials can be detected. Two examples of this plot type can be seen in figure 3.3[†]

[†]Source code for the layered feature plots can be found in appendix A.7 and the results are in appendix A.8

From the layered feature plots it is confirmed that many features have some trials that seems to deviate from the normal pattern of the feature. Using feature P6 as an example, it is seen in figure 3.3a that P6 in almost all trials lies within the range [0,2000]. But one trial have a max value larger than 200000. This behaviour was hinted in section 3.2 and now there is a visual confirmation.

The question is, how should this “outlier” trial of feature P6 be handled? One approach is to exclude the trial from the dataset. But as nothing about the origin of the dataset is known, it can’t be justified with a common sense argument, and if this trial is excluded then how is the line between outlier and normal drawn? There is a couple of other trials where the P6 feature has values as high as 20000, but should this be seen as outliers? And are there other ways a trial could be an outlier than the scale? In figure 3.3b most trials have the same structure, but the trial with the lowest mean has a noticeable different shape than the other trials. It somehow has a baseline at the value 400 and then oscillates a higher frequency than the other trials. Is this an outlier? If it is regarded as an outlier many other trials could possibly be classified as outliers, and we’re are still only looking at one feature. If the same amount of cleaning is done for the other features, not much is left for the “normal” dataset.

The immediate conclusion is that the outlier detection easily becomes, either too all-encompassing or is done in an inconsistent manner. An alternative way to tackle the trials with datapoints with large values, could be to do some kind of variable transformation (eg. taking the log) of the feature. This could be a solution and the idea is investigated a little more in section 5.

3.4.3 Features that are constant in a trial

Another type of potential outlier trials, are all the trials where one or more features are constant throughout the trial. This could be a sign of a failure in the measurement equipment, and then maybe the trial should be regarded as an outlier? A great example is the IsAlert feature, that is the feature that should be predicted by the classifier. In 142 out of the 400 trials the IsAlert feature is constant and in 127 of the trials with a constant

value, the driver isn't alert. At first it may seem a bit odd, and indeed some of the contestants of The Ford Challenge revealed that they had excluded all trials where the IsAlert feature was 0 throughout the trial (see [David \(2011\)](#)). But it is speculative to do, since the precise meaning of the IsAlert feature isn't known, let alone the fact that, measuring human alertness with a binary variable is speculative by itself.

As it seems to be speculative to exclude features with constant values, this idea won't be pursued any further. Instead the focus changes from "finding features to get rid of" till "finding features that could be beneficial for classification" as the next section is all about finding features that can discriminate between alert/not-alert drivers.

3.5 Finding possible discriminating features

In this section all features are examined in an attempt to find one or more features that are in some way correlated with the IsAlert feature that should be predicted.

3.5.1 Testing binary features

A tempting first attempt at finding a feature that can be used to predict the IsAlert feature, is to look at the two binary features V5 and E9 and see if they are correlated with the predictor **TODO: What is the IsAlert feature called?**.

3.5.2 Scatterplots

A great way to visualize whether a pair of features can be used to discriminate between datapoints where the driver is alert and where he/she is not alert, is to create scatterplots of the possible pair of features. Since feature V7, V9 and P8 are removed and feature E3, E7, E8, E9, V3, V5 and V10 was found to be categorical, only 20 features are left for the scatterplots. Since

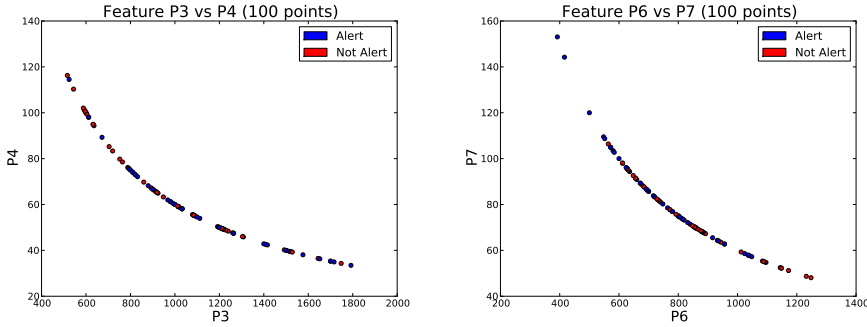


Figure 3.4: Scatterplots of the feature pairs (P3, P4) and (P6, P7). An arithmetic relationship between the features is clearly present

the order of the two features in a scatterplot doesn't matter there is a total of

$$\binom{20}{2} = 190 \text{ combinations}$$

All combinations of the 20 features was plotted[‡], but no interesting results was found; at least not regarding feature pairs that could classify the IsAlert feature. One interesting result was found though. As can be seen in figure ?? the features pairs (P3, P4) and (P6, P7) both appear to have an arithmetic relationship. A quick calculation shows that P3 multiplied by P4 is 60000 ± 0.2 for all rows in the dataset, and likewise for P6 and P7. We can therefore ignore one of the features from each pair. **TODO: Am I sure? Details can't hide in the decimal difference?**

3.6 Principal Component Analysis

3.6.1 Theory[§]

The idea in Principal Component Analysis (PCA) is to find a mapping from a d -dimensional input space, to a new k -dimensional space (with

[‡]Source code can be found in appendix A.9 and the results in appendix A.10

[§]Theory section is based on Alpaydin (2010, p.113)

$k < d$), while keeping as much of the variation in the input data, as possible. To begin with let $k = 1$. Then the subspace is identified by a d -dimensional vector \mathbf{w}_1 , that without loss of generality is assumed to have length $\|\mathbf{w}_1\| = 1$.

For any vector \mathbf{x} in the input space, the projection onto \mathbf{w}_1 is given by

$$z_1 = \mathbf{w}_1^T \mathbf{x}$$

and we now wish to maximize the variance of z_1 . By the definition of variance we get (with $\boldsymbol{\mu} = E[\mathbf{x}]$)

$$\begin{aligned} \text{Var}[z_1] &= \text{Var}[\mathbf{w}_1^T \mathbf{x}] \\ &= E[(\mathbf{w}_1^T \mathbf{x} - \mathbf{w}_1^T \boldsymbol{\mu})^2] \\ &= E[(\mathbf{w}_1^T \mathbf{x} - \mathbf{w}_1^T \boldsymbol{\mu})(\mathbf{w}_1^T \mathbf{x} - \mathbf{w}_1^T \boldsymbol{\mu})] \\ &= E[\mathbf{w}_1^T (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T \mathbf{w}_1] \\ &= \mathbf{w}_1^T E[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T] \mathbf{w}_1 \\ &= \mathbf{w}_1^T \Sigma \mathbf{w}_1 \end{aligned}$$

where $\Sigma = \text{Cov}[\mathbf{x}]$. So to maximize the variance of z_1 is the same as maximizing $\mathbf{w}_1^T \Sigma \mathbf{w}_1$, subject to the constraint that $\mathbf{w}_1^T \mathbf{w}_1 = 1$. By introducing a Lagrange multiplier α , we get

$$\max_{\mathbf{w}_1} \mathbf{w}_1^T \Sigma \mathbf{w}_1 - \alpha (\mathbf{w}_1^T \mathbf{w}_1 - 1)$$

This is a normal, unconstrained maximization problem. Taking the derivative with respect to \mathbf{w}_1 , and setting this equal to 0, gives

$$\begin{aligned} 2\Sigma \mathbf{w}_1 - 2\alpha \mathbf{w}_1 &= 0 & \Leftrightarrow \\ \Sigma \mathbf{w}_1 &= \alpha \mathbf{w}_1 \end{aligned}$$

which is true exactly when \mathbf{w}_1 is an eigenvector of Σ and α is the corresponding eigenvalue. Giver det mening på binære variable?

3.7 Conclusions

Part II

Modelling

CHAPTER 4

Theory of classification

In this chapter some of the theory of classification is explained. The discussion is constrained to binary classification, of which the Ford Challenge classifier, is an example. First a general problem definition is given and notation is introduced. Then three different approaches to classification are introduced, and some pros and cons of each approach are mentioned. After this two concrete examples of classification models are introduced (Logistic Regression and Feed Forward Neural Network), and they are related to the three general approaches of classification. Finally different ways to measure the performance of different classifiers are discussed, including AUC that was used as the grading method in The Ford Challenge.

4.1 The binary classification problem^{*}

In a binary classification problem a binary outcome t must be predicted from a d -dimensional input vector $\mathbf{x} = [x_1, x_2, \dots, x_d]^T$. The input vector

^{*}This section is loosely based on [Wasserman \(2004, sec 22.1-22.2\)](#)

represents an event, and the outcome variable t represents the assignment of the event to one of two classes.

EXAMPLE 4.1 *In The Ford Challenge the input is an instant in a driving situation, and this input should be assigned to either the class alert, or the class not-alert. The input is represented by an 30-dimensional vector and the assignment to a class is represented by $t = 0$ meaning not-alert and $t = 1$ alert.*

To make the prediction possible, a trainingset \mathcal{T} is given consisting of n pairs of input vectors and corresponding, known outcomes.

$$\mathcal{T} = \{(\mathbf{x}_i, t_i)\}$$

Based on this trainingset, a classification rule f is learned, that can predict the outcome, of yet unknown input vectors. Therefore a classification rule is a function $f : \mathbb{R}^d \rightarrow \{0, 1\}$, that is used to make the prediction $t = f(\mathbf{x})$ on a new input.

A common way to obtain a classification rule is by the Bayes classification rule. First define $p_k(\mathbf{x})$ as

$$p_k(\mathbf{x}) = P(t = k | \mathbf{x}), \quad k \in \{0, 1\}$$

and then the Bayes classification rule is defined by

DEFINITION 4.1 *The Bayes classification rule f^* is given by*

$$f^*(\mathbf{x}) = \begin{cases} 1 & \text{if } p_1(\mathbf{x}) > p_0(\mathbf{x}) \\ 0 & \text{else} \end{cases}$$

Since $p_0(\mathbf{x}) = 1 - p_1(\mathbf{x})$, the Bayes classification rule can also be written

$$f^*(\mathbf{x}) = \begin{cases} 1 & \text{if } p_1(\mathbf{x}) > \frac{1}{2} \\ 0 & \text{else} \end{cases}$$

By the Bayes classification rule we have got a theoretical classification rule. In practice the posterior class probabilities $p_0(\mathbf{x})$ and $p_1(\mathbf{x})$ are unknown, so the trainingset must be used to find some approximation for these probabilities. The problem of approximating a probability distribution from a given data set is a classic statistical problem. The next section gives one way to solve this problem.

4.2 Parametric models and maximum likelihood

Two distinct ways to approximate the probability distribution $P(t|\mathbf{x})$ is using either a parametric or a non-parametric approach. The focus in this report is on the parametric approach. The distribution $P(t|\mathbf{x})$, is therefore approximated by a distribution $\hat{P}(t|\mathbf{x})$ restricted to a class of distributions

$$\mathcal{F} = \left\{ f(t|\mathbf{x}, \mathbf{w}) \right\}$$

where the distributions $f_{\mathbf{w}} \in \mathcal{F}$ are uniquely identified by their parameter $\mathbf{w} = [w_1, w_2, \dots, w_k]^T$. By restricting the approximating distribution $\hat{P}(t|\mathbf{x})$ to the set \mathcal{F} , the problem of approximating $P(t|\mathbf{x})$ reduces to the problem of estimating the parameter \mathbf{w} .

4.2.1 Maximum likelihood

There exists some different techniques to estimate the parameter \mathbf{w} in a parametric model.[†] In this section we look at the most common method for estimation of the parameter in a parametric model, namely the maximum likelihood method.

The approximation $\hat{P}(t|\mathbf{x})$ of the posterior class probability is restricted to the class \mathcal{F} . Using the trainingset \mathcal{T} , the likelihood function is now defined as the probability of obtaining the data in the trainingset, given as a function of the parameter \mathbf{w} . More formally[‡]:

DEFINITION 4.2 *The likelihood function is defined by*

$$\mathcal{L}(\mathbf{w}) = \prod_{i=1}^n f(t_i|\mathbf{x}_i, \mathbf{w}), \quad (t_i, \mathbf{x}_i) \in \mathcal{T}$$

and the log-likelihood function is defined by $\ell(\mathbf{w}) = \log \mathcal{L}(\mathbf{w})$.

It is worth noticing that the definition of the likelihood function assumes that the t_i 's are independent random variables. An assumption that in the

[†]See eg. Wasserman (2004, Sec.9)

[‡]Taken directly from Wasserman (2004, p.122)

case of The Ford Challenge may be a bit questionable. **TODO: Elaborate this point.**

Given the definition of the likelihood function, the maximum likelihood estimator is now defined by

DEFINITION 4.3 *The maximum likelihood estimator $\hat{\mathbf{w}}$ is the value of \mathbf{w} that maximizes $\mathcal{L}(\mathbf{w})$.*

Notice that as the logarithm is a monotonic increasing function, the maximum of $\mathcal{L}(\mathbf{w})$ is equal to the maximum of $\ell(\mathbf{w})$. This turns out to be handy, since the log-likelihood often is easier to work with, than the likelihood function.

With these definitions it is now time to introduce two concrete examples of classification methods.

4.3 Logistic Regression

In the logistic regression the parametric form of the posterior class probability $P(t = 0|\mathbf{x})$ is assumed to be

$$P(t = 0|\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + w_0)}} \quad (4.1)$$

where $\mathbf{w} = [w_1, w_2, \dots, w_d]^T$. The expression may seem a bit random, but it comes from the fact that (with $p_0 = P(t = 0|\mathbf{x})$)

$$p_0 = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + w_0)}}$$

can be written

$$\begin{aligned} p_0 &= \frac{e^{\mathbf{w}^T \mathbf{x} + w_0}}{1 + e^{\mathbf{w}^T \mathbf{x} + w_0}} && \Leftrightarrow \\ p_0 + p_0 e^{\mathbf{w}^T \mathbf{x} + w_0} &= e^{\mathbf{w}^T \mathbf{x} + w_0} && \Leftrightarrow \\ \frac{p_0}{1 - p_0} &= e^{\mathbf{w}^T \mathbf{x} + w_0} && \Leftrightarrow \\ \log \frac{p_0}{1 - p_0} &= \mathbf{w}^T \mathbf{x} + w_0 \end{aligned}$$

And therefore the parametric form in (4.1), comes from the assumptions that the log-odds of p_0 is linear in the input vector \mathbf{x} .

Now to find an expression for the likelihood function, the probability of getting the data in the trainingset must be expressed. With $p_i = P(t = 0|\mathbf{x}_i)$ the likelihood function can be written as **TODO: Mention the Bernoulli distribution**

$$\mathcal{L}(\mathbf{w}) = \prod_{i=1}^n p_i^{t_i} (1 - p_i)^{1-t_i}$$

which gives the log-likelihood function

$$\ell(\mathbf{w}) = \sum_{i=1}^n t_i \log p_i + (1 - t_i) \log(1 - p_i)$$

The maximum of the log-likelihood function gives the maximum likelihood estimator of \mathbf{w} , so the derivative of ℓ wrt. \mathbf{w} need to be derived. Using the fact that

$$\frac{\partial p_i(\mathbf{w})}{\partial \mathbf{w}} = p_i(1 - p_i)\mathbf{x}_i$$

gives that

$$\begin{aligned} \frac{\partial \ell(\mathbf{w})}{\partial \mathbf{w}} &= \sum_{i=1}^n t_i \frac{1}{p_i} p_i(1 - p_i)\mathbf{x}_i + (1 - t_i) \frac{1}{1 - p_i} (-p_i)(1 - p_i)\mathbf{x}_i \\ &= \sum_{i=1}^n t_i(1 - p_i)\mathbf{x}_i - (1 - t_i)p_i\mathbf{x}_i \\ &= \sum_{i=1}^n (t_i - p_i)\mathbf{x}_i \end{aligned}$$

The maximum likelihood estimator can now be found by solving $\frac{\partial \ell(\mathbf{w})}{\partial \mathbf{w}} = 0$, but since p_i isn't linear wrt. \mathbf{w} the equation must be solved numerically. An effective algorithm called iterative reweighted least squares can solve this problem (Bishop, 2006, p.207). **TODO: Show that the Hessian is Negative Semi Definite and therefore an unique maximum exists**

4.3.1 The decision boundary

An important property of a classification method is how complicated data sets the method is able to separate in the appropriate classes. The decision

boundary of a classification method gives some information about how well it separates input data. The decision boundary is the set of points in the input space that lays at the boundary between input points that are classified as class 0 and input points that classify as class 1.

EXAMPLE 4.2 *A trainingset consists of n d -dimensional input data $\mathbf{x}_i = [x_{1i}, x_{2i}, \dots, x_{di}]^T$ and corresponding classes t_i . Using linear regression, approximations of the posterior class probabilities $\hat{p}_k(\mathbf{x}) = \hat{P}(t = k|\mathbf{x})$, $k \in \{0, 1\}$ are given. Using these with the Bayes classification rule gives*

$$f(\mathbf{x}) = \begin{cases} 1 & \text{if } \hat{p}_1(\mathbf{x}) > \hat{p}_0(\mathbf{x}) \\ 0 & \text{else} \end{cases}$$

The decision boundary is given by all \mathbf{x} where $\hat{p}_0(\mathbf{x}) = \hat{p}_1(\mathbf{x})$. But by the definition of the linear regression

$$\log \frac{\hat{p}_1(\mathbf{x})}{\hat{p}_0(\mathbf{x})} = \mathbf{w}^T \mathbf{x} + w_0$$

So for all points at the decision boundary the following holds

$$\mathbf{w}^T \mathbf{x} + w_0 = 0$$

This shows that the decision boundary of the logistic regression is linear in the inputs.

The classification methods with linear decision boundaries, defines a family of classification methods with simple analytical and computational properties (Bishop, 2006, p.179). This simplicity comes at the cost of the models having a limited capability to separate complex data sets.

4.3.2 Model complexity

Although the logistic regression has a limited complexity, its ability to separate complex datasets can be increased by introducing feature transformations. Instead of just fitting the logistic regression to the d -dimensional input $\mathbf{x} = [x_1, x_2, \dots, x_d]$, the squared transformations of the elements x_i are added to the input. This gives the $2d$ -dimensional input

$$\mathbf{x}_e = [x_1, \dots, x_d, x_1^2, \dots, x_d^2]$$

This gives nonlinear decision boundaries in the original input space, which increases the number of datasets that can be separated. Notice that the decision boundary is still linear in the extended $2d$ -dimensional inputspace. Instead of introducing transformed, non-linear features in the logistic regression, other methods exist, that have the non-linearity build-in. One example of such a model is the Neural Network.

4.4 Feed Forward Neural Network

The Neural Network originated from an attempt to create a mathematical model of information processing in the human brain (Bishop, 2006, p.226). But the Neural Network turned out to be useful in many practical applications, and it is now seen as a standard statistical model (Hastie et al., 2009, p.392).

4.4.1 Model complexity

4.4.2 Regularization

4.5 Measuring classifier performance

When a classifier has been trained, an estimate of the classifiers expected performance on future inputs, should be calculated. To calculate this estimate, it needs to be specified exactly what is meant by a classifiers performance. For that a loss function is defined

DEFINITION 4.4 *Let t_0 be the true class of a corresponding input \mathbf{x}_0 , and f be a classification rule. The loss function*

$$L(t_0, f(\mathbf{x}_0))$$

then gives the cost of making the prediction $f(\mathbf{x})$, when the true class is t_0 .

The common way to define the loss function is by

$$L(t, y) = \begin{cases} 1 & \text{if } t \neq y \\ 0 & \text{else} \end{cases} \quad (4.2)$$

This assumes that the cost of predicting class 0 when the truth is class 1, and visa versa is the same. In many applications this may be questionable. In The Ford Challenge the cost of predicting alert, when the truth is not-alert, might be a more severe error, than the other way around; although it depends on what the Ford classifier is used for. In the rest of this section, the loss function given in (4.2) is assumed.

When a loss function has been defined, what would be interesting to know, is the expected loss on future data for a classification rule f .

$$\text{Err} = \mathbb{E}[L(t, f(\mathbf{x}))]$$

where t and \mathbf{x} is random variables and the expectation is taken on their joint distribution. The expected loss isn't normally known so it needs to be estimated somehow. For that the average loss over a dataset for a given classification rule can be used. For an arbitrary dataset $\mathcal{D} = \{(\mathbf{x}_i, t_i)\}, i = 1 \dots n$, classification rule f and loss function L , the average loss is calculated by

$$\overline{\text{err}} = \frac{1}{n} \sum_{i=1}^n L(t_i, f(\mathbf{x}_i))$$

4.5.1 Splitting the dataset

It is extremely important that the dataset used to calculate the average loss, isn't in anyway overlapping with the dataset that was used to train the classifier. If there is a large dataset available for a classification task, the dataset should ideally be split into three separate datasets[§]

- The training set \mathcal{T} - This dataset is used to train models. More specifically the training set is used to find the maximum likelihood estimator, for a parametric model

[§]Loosely based on [Hastie et al. \(2009, p.222\)](#)

- The validation set \mathcal{V} - This dataset is used for model selection. As mentioned in section 4.3.2, the complexity of classification methods can be altered to make the method fit the training data better. The validation set gives an estimate of how well models of different complexity levels, predicts future datapoints.
- The publication set \mathcal{P} - This dataset is used, when the final model has been selected, and an estimate of the performance of this model.

There is not a general rule for how many percent of the original dataset that should be assigned to each of the three datasets. In Hastie et al. (2009, p.222), a split of 50% for training and 25% for both validation and publication, is suggested.

EXAMPLE 4.3 *In The Ford Challenge a dataset of 500 trials is given. Following the suggestion by Hastie et al. (2009), the dataset is split in a training set of 250 trials and a validation set and publication set of both 125 trials. It is then decided to train a logistic regression model to classify drivers. The complexity of the model is increased by including various feature transformations. The average loss of the different models is calculated on the validation set, and the models with the lowest average loss is selected. Finally the average loss over the publication loss is used as an independent estimate of the expected loss on future data.*

The above method of making three distinct datasets out of the original, is only possible, when the original dataset is large. And even with a large dataset, it may seem a bit wasteful to only have tested the final model against 50% of the original dataset **TODO: Does this make sense?**. A way to circumvent this problem is by using cross validation.

4.5.2 Cross validation

Cross validation is a technique to squeeze out more information of a dataset, than by splitting the dataset in three parts as described in section 4.5.1 **TODO: Is this right? I have troubles defining cross validation? Is it only a technique for model selection? When assessing performance on**

test/publication set cross validation could be used? Recommendation of 5- or 10-fold cross validation (Hastie et al., 2009, p.243)

4.5.3 AUC

Mentioning critiques of AUC?

4.6 Model selection

4.6.1 Over- and underfitting

4.6.2 Forward feature selection

4.6.3 Regularization

CHAPTER 5

Recreating winning approach

This is the first chapter in the modelling part. In this chapter, the classification method of the Ford Competition winner will be attempted to be recreated. The goal is to get a result – as measured by the AUC – that is close to the winners approach.

5.1 The winning approach

The competition winner has described his approach in [Inference \(2011a\)](#), which will just be called “the paper” in the rest of this section. The main points in the paper will be described in this section.

In the paper two key observations are mentioned. First of all, the winner observed that in most of the trials, the driver was either alert most of the time, or not alert for most of the time. Secondly the winner observed that there was a mismatch between the AUC-scores achieved on the trainingset and on the testset. Since it wasn’t possible to get access

to the testset of the competition, the second observation isn't of any direct use for this report. The two observations though, gives an explanation of why the winner chose the model he did.

The fact that most trials was either alert-trials or not-alert-trials, the winner thought it would make sense to calculate aggregated features within trials. He therefore included the mean and the standard deviation of all features, aggregated backwards, within a trial. As a standard the mean features was prefixed "m" and the standard deviation "sd".

EXAMPLE 5.1 *In the k 'th observation of trial t , the value of feature $mP6$, is the mean of feature $P6$, for observation $0, 1, 2, \dots, k - 1$ in trial t*

The fact that there was a mismatch between the trainingset and testset of the competition, made the winner focus on simple models, since he expected that the trainingset wasn't capturing all aspects of the testset, and extrapolation was therefore needed.

5.1.1 Classification method

The method used by the winner was a logistic regression, trained on only the three features $sdE5$, $V11$ and $E9$. The three features was selected "... based on diagnostics of the logistic regression ..." (Inference, 2011a, p.3). The final model was a logistic regression with the following weights

$$\log \frac{P(t = 0|\mathbf{x})}{P(t = 1|\mathbf{x})} = -392.4317 \cdot sdE5 + 0.2209 \cdot V11 + 3.6544 \cdot E9$$

The winner do not mention any intercept parameter, although it seems unlikely that no intercept was included in the model (Mørup, 2011). Based on the above parameters, the winner achieved an AUC-score of 0.8492 on the testset. As mentioned in the previous section, the AUC-score of the testset differed from the AUC-score on the trainingset. In figure 2 in Inference (2011a) it is seen that an AUC-score of 0.8492 on the testset, corresponded to a AUC-score of 0.82 ± 0.01 on the trainingset. Since only the trainingset is available for this report, the goal is to achieve an AUC-score of approximately 0.82 ± 0.01 .




Figure 5.1: Reproduction of figure 2 in *Inference (2011a)*. The models that achieves a Test AUC higher than 0.8492 are using future observations to predict the IsAlert feature, which wasn't allowed in the competition. See *Inference (2011a)* for more information.

Run	sdE5	V11	E9	Intercept	AUC
1	-87.2063	0.2239	3.6578	-5.3642	0.8051
2	-353.2778	0.2157	3.6969	-4.9869	0.8185
3	-87.4096	0.2246	3.6644	-5.3763	0.8048
4	-352.1050	0.2159	3.6960	-4.9905	0.8208
5	-86.9788	0.2231	3.6455	-5.3440	0.8097
6	-87.1768	0.2239	3.6487	-5.3539	0.8060
7	-87.1603	0.2244	3.6575	-5.3666	0.8074
8	-87.3095	0.2245	3.6618	-5.3736	0.8078
9	-353.1568	0.2150	3.6951	-4.9778	0.8177
10	-87.3047	0.2245	3.6461	-5.3590	0.8095

Table 5.1: Results of running a 10-fold cross validation logistic regression, on the features used by the winner.

5.2 Recreating the winning approach

It is now time to recreate the logistic regression of the winner. A 10-fold cross validation method is used. The trainingset is therefore split in 10 parts, and one part is repeatedly used as validationset while the other 9 parts, are used to train the logistic regression. This gives 10 estimates of the true AUC-score, and 10 different parameter vectors. The results can be seen in table 5.1 and the source code used to calculate the results can be found in appendix A.11

From the results table it is seen that the single best run is the fourth run, that gives the model

$$\log \frac{P(t=0|\mathbf{x})}{P(t=1|\mathbf{x})} = -352.1050 \cdot \text{sdE5} + 0.2159 \cdot \text{V11} + 3.6960 \cdot \text{E9} - 4.9905$$

This model achieves an AUC-score of 0.8208 on the validationset. From the results, it is also seen that there are two distinct patterns in the parameters, that are given by the logistic regression. On 7 out of the 10 runs the weight of the sdE5 feature is about 87 and on the rest of the runs it is 352-353. The runs with a sdE5 weight around 352-353 are consistently giving better AUC-scores than the other runs. The three runs with the sdE5 feature of 352-353, all have feature weights that is close to the weights used in the winning model. The AUC-scores achieved by these three models also comes close to the AUC-scores of the winner.

5.2.1 Statistics on the AUC score

From the results of the 10-fold cross validation logistic regression, 10 AUC scores was obtained. If it is assumed that the 10 AUC scores are a sample from the population of all AUC scores on unknown datasets, a confidence interval for the AUC score can be calculated. The mean $\hat{\mu}$ and standard deviation $\hat{\sigma}$ of the AUC score are calculated as

$$\hat{\mu} = 0.8107 \quad \text{and} \quad \hat{\sigma} = 0.005980$$

TODO: Calculate confidence interval for AUC estimate

CHAPTER 6

Improving the winning approach

In this chapter it is examined whether the winning approach can be improved or not. The classification method will still be logistic regression, but it is tried to improve the results, by doing feature selection across all features. Two feature selection techniques are tried. First a forward selection is performed, on a subset of the whole trainingset, and then a L^1 regularization. Finally the performance is tried to be improved by substituting the running features, by other features.

6.1 Forward selection

A forward selection is now performed to find the features that gives the best classification. The logistic regression model starts out with no features, and then the feature that improves the performance the most is added to the regression model. This is repeated until the performance is no longer increased. The features, that the forward selection chooses be-

Feature added	AUC
V11	0.6978
E9	0.8089
sdE1	0.8333
mE6	0.8500
sdE4	0.8539
⋮	⋮
mP1	0.8858
sdE5	0.8857

Table 6.1: The first five and the last two features added in the forward selection.

tween are all the features; both the original, and the running mean and running standard deviation. This gives a total of 90 features to choose between. Since all candidates for features to add to the model, must be tested in every iteration, a worst case scenario is that a total of

$$\sum_{i=1}^{90} i = \frac{90 * 89}{2} = 4005$$

logistic regression fittings must be done. Since a standard library (scikit-learn) is used for the training, the training dataset is restricted to 10000 rows and the validation dataset to 5000 rows. And only a simple hold-out cross validation is performed. The script for the forward selection can be seen in appendix A.12 and all the results are in appendix A.13. The first five selected features and the last two are shown in table ?? A total of 48 features out of the 90 features was selected. This is more features than expected**TODO: Is it?** and the AUC that is achieved, when all 48 features are in the model is quite high (0.8857). The 10000 datarows for the trainingset and the 5000 datarows of the validationset was chosen uniformly across the whole dataset**TODO: Find good names for the different datasets**, so there is no reason to believe that any connection between the trainingset and the validationset, makes the AUC score artificially high. It is worth noticing that the top two features (V11, E9) are also used in the To get a better idea of the precise performance of the model with the 48 selected features, a 10-fold cross validation on the original testset is now performed.

Run	AUC
1	0.9290
2	0.9287
3	0.9297
4	0.9311
5	0.9348
6	0.9332
7	0.9332
8	0.9286
9	0.9327
10	0.9344

Table 6.2: Results from running a 10-fold cross validation on the report testset, with the full 48 feature logistic regression model, selected by forward selection

6.1.1 Testing performance on the original testset

By forward selection a model with 48 features have been selected. The AUC score achieved when the 48'th feature was added was 0.8857. But this score was achieved on a small validationset of 5000 rows. Now it is tested how the model performs on the original publicationset. Therefor a 10-fold cross validation is performed on the publicationset. The source code is almost identical with the source code used to recreate the winning approach, and this can be found in appendix A.11. The results can be seen in table ??, and they are a bit surprising.

The AUC score of the 10 runs falls in the interval [0.928,0.935], which is significantly above the 0.8857**TODO: Morten: God forklaring? Måske forskellen på de to datasæts størrelse 10000/5000 vs. 109800/12200** that was achieved when running the forward selection (see table ??). It would have been great if the model could have been tested on the original Ford testset, since there have been some talk in the Kaggle forum (?) about models, achieving AUC scores above 0.9, that still scored below 0.8 on the ford test set. Since it isn't possible to test the model on the ford test-set, it can only be concluded that on the report testset, this model scored auc scores around 0.93**TODO: Statistics.**

With a model with an AUC score of 0.93, the goal of improvement of the winners model, has been achieved. But this improvement was created

Run	AUC
1	0.8335
2	0.8373
3	0.8382
4	0.8326
5	0.8373
6	0.8339
7	0.8329
8	0.8325
9	0.8326
10	0.8318

Table 6.3: Results from running a 10-fold cross validation on the report testset, with the top 3 features selected by forward selection

by a much more complex model than the winning model, and one of the reasons the winning model was kept simple, was exactly the problem with models scoring high on validationsets, but low on the ford testset. To make a (maybe) more fair comparison, the performance of a model with the top 3 features from the feature selection is calculated. The procedure is exactly as with all 48 features and the new results are seen in table 6.3

TODO: IMPORTANT: I have tested the forward selection models on the testset, but the original winner model on the trainingset. If I train the original on the testset I can't recreate his parameter weights. If I train the forward selection models on the trainingset I use the same dataset for feature selection and validation

TODO: To compare the performance of the top 3 forward selected features vs. the winning model, a statistical procedure would be great. There seem to be an improvement in the forward selected model, but it isn't clear cut.

CHAPTER 7

Other classification methods

Here I describe some alternatives to the logistic regression used by Inference. Hoping to get a result or two from SVM or Neural Network.

Part III

Workflow and discussion

CHAPTER 8

Workflow and tools

Describing software used, workflow using github, writing sessions, evaluate my performance, mention the small improvement of roccurve method.

CHAPTER 9

Discussion

Bla, bla, bla.

CHAPTER 10

Conclusion

More bla, bla, bla.

CHAPTER **A**

Appendices

A little introduction and then a new page

A.1 Calculating common statistics - Code

```
# sessions/9-data-exploration/src/calculate_summary_statistics.py

from json import dump

from src.data_interface import trd, L
from src.utils import get_path

sess_root = get_path(__file__) + '/../scr'

summary_statistics = ['min', 'max', 'mean', 'std']
features_to_calculate = L[2:]
calculations = {}

for feature_name in features_to_calculate:
    calculations[feature_name] = {}
    for statistic in summary_statistics:
        stat_function = getattr(trd.get_feature(feature_name), statistic)
        calculations[feature_name][statistic] = stat_function()

f = open(sess_root + '/summary_statistics.json', 'w')
dump(calculations, f, indent=4)
f.close()
```

Code Listing A.1: https://github.com/alphabits/ford_challenge/tree/master/sessions/9-data-exploration/src/calculate_summary_statistics.py

A.2 Calculating common statistics - Result

Feature	Mean	Min	Max	Std.Dev
E1	10.54	0.00	243.99	14.00
E10	63.38	0.00	127.00	19.23
E11	1.37	0.00	52.40	5.37
E2	105.05	0.00	360.00	128.33
E3	0.30	0.00	4.00	1.03
E4	-3.99	-250.00	260.00	34.80
E5	0.02	0.01	0.02	0.00
E6	358.50	260.00	513.00	27.84
E7	1.81	0.00	25.00	2.95
E8	1.39	0.00	9.00	1.62
E9	0.87	0.00	1.00	0.33
IsAlert	0.58	0.00	1.00	0.49
P1	35.45	-22.48	101.35	7.36
P2	12.01	-45.63	71.17	3.74
P3	1028.00	504.00	2512.00	309.70
P4	63.99	23.89	119.05	19.76
P5	0.18	0.04	27.20	0.39
P6	843.73	128.00	228812.00	2795.32
P7	78.40	0.26	468.75	18.79
P8	0.00	0.00	0.00	0.00
V1	75.58	0.00	129.70	44.94
V10	3.28	1.00	7.00	1.27
V11	11.59	1.68	262.45	8.43
V2	-0.04	-4.79	3.99	0.42
V3	569.78	240.00	1023.00	299.02
V4	21.24	0.00	484.49	67.21
V5	0.18	0.00	1.00	0.38
V6	1699.34	0.00	4892.00	626.27
V7	0.00	0.00	0.00	0.00
V8	12.46	0.00	82.10	11.54
V9	0.00	0.00	0.00	0.00

A.3 Calculate unique values of features - Code

```
from json import dump

from src.data_interface import trd, L
from src.utils import get_path

sess_root = get_path(__file__) + '/../src'

features_to_calculate = L[2:]
unique_values = {}

for feature_name in features_to_calculate:
    unique_values[feature_name] = trd.get_feature(feature_name).unique_values()

f = open(sess_root + '/unique_values.json', 'w')
dump(unique_values, f, indent=4)
f.close()
```

Code Listing A.2: https://github.com/alphabits/ford_challenge/tree/master/sessions/9-data-exploration/src/calculate_unique_values.py

```
from json import dump
import numpy as np

from src.data_interface import trd, L
from src.utils import get_path

sess_root = get_path(__file__) + '/../'

features_to_calculate = L[2:]
trials = list(trd.trial_id_list)
calculations = {}

for feature_name in features_to_calculate:
    tmp = {"trial_results": {}}
    for trial_id in trials:
        unique_values = np.unique(
            trd.get_trial(trial_id).get_feature(feature_name).view())
        tmp["trial_results"][trial_id] = unique_values.size
    tmp["max"] = max(tmp["trial_results"].values())
    tmp["min"] = min(tmp["trial_results"].values())
    calculations[feature_name] = tmp

f = open(sess_root + '/src/unique_values_pr_trial.json', 'w')
dump(calculations, f, indent=4)
f.close()
```

Code Listing A.3: https://github.com/alphabits/ford_challenge/tree/master/sessions/9-data-exploration/src/calculate_unique_values_pr_trial.py

```
import json

from src.utils import get_path

path = get_path(__file__)

def get_dict(file_name):
    f = open(path + '/' + file_name, 'r')
    d = json.load(f)
    f.close()
    return d

unique_pr_trial = get_dict('unique_values_pr_trial.json')
unique_all_data = get_dict('unique_values.json')

values_combined = {}

for k in unique_all_data:
    tmp = {}
    tmp["all_data"] = unique_all_data[k]
    tmp["min_pr_trial"] = unique_pr_trial[k]["min"]
    tmp["max_pr_trial"] = unique_pr_trial[k]["max"]
    values_combined[k] = tmp

f = open(path + '/unique_values_combined.json', 'w')
json.dump(values_combined, f, indent=4)
f.close()
```

Code Listing A.4: https://github.com/alphabits/ford_challenge/tree/master/sessions/9-data-exploration/src/unique_values_combined.py

A.4 Calculate unique values of features - Result

Feature	Min in trial	Max in trial	Whole dataset
E1	1	131	12265
E10	1	86	121
E11	1	79	116
E2	1	129	28502
E3	1	3	3
E4	1	97	249
E5	1	134	254
E6	1	116	247
E7	1	15	25
E8	1	10	10
E9	1	2	2
IsAlert	1	2	2
P1	967	1208	105715
P2	1155	1207	128666
P3	24	106	406
P4	24	106	406
P5	34	114	1070
P6	3	99	406
P7	3	99	406
P8	1	1	1
V1	1	969	12374
V10	1	5	5
V11	875	1206	166455
V2	1	76	90
V3	1	29	33
V4	1	204	324
V5	1	2	2
V6	1	775	2727
V7	1	1	1
V8	1	239	323
V9	1	1	1

A.5 Creating boxplots - Code

TODO: Include source code

A.6 Creating boxplots - Result

TODO: Include boxplots

A.7 Creating layered feature plots - Code

A.8 Creating layered feature plots - Result

A.9 Scatterplots - Code

```
import itertools as it

import matplotlib.pyplot as plt
from matplotlib.patches import Rectangle
import numpy as np

from src.dataloaders import trainingset
from src.utils2 import get_path, c, L

L = list(L[4:])
D = trainingset()

path = get_path(__file__) + '/../'
savepath_template = '{0}/plots/scatterplots/{1}-{2}.pdf'
num_points = 100

rows = np.random.random_integers(0,D.shape[0]-1, num_points)
data = D[rows,:]

colors = map(lambda x: 'blue' if x==1 else 'red', data[:,c('IsAlert')])
blue = Rectangle((0,0),1,1,fc='b')
red = Rectangle((0,0),1,1,fc='r')

exclude = ['V7', 'V9', 'P8', 'E3', 'E7', 'E8', 'E9', 'V3', 'V5', 'V10']
features = [f for f in L if f not in exclude]

for f1, f2 in it.combinations(features, 2):
    #for f1, f2 in [['E4', 'E5']]:
        plt.title('Feature {0} vs {1} ({2} points)'.format(f1, f2, num_points),
                  {'size': 20})
        plt.legend((blue, red), ('Alert', 'Not Alert'))
        plt.scatter(data[:,c(f1)], data[:,c(f2)], c=colors)
        plt.gca().set_xlabel(f1, {'size': 18})
        plt.gca().set_ylabel(f2, {'size': 18})
```

```
plt.savefig(savepath_template.format(path,f1,f2), format='pdf',
            papertype='a4')
plt.cla()
```

Code Listing A.5: https://github.com/alphabits/ford_challenge/tree/master/sessions/24-writing-helper-scripts/scripts/rewrite-scatterplot-script.py

A.10 Scatterplots - Result

A.11 Recreating the winning approach - Code

```
from __future__ import division

import json
import random

import numpy as np
import matplotlib.pyplot as plt

from src.utils2 import c_ex as c, get_path
import src.dataloaders as d
from src.logistic import fit_logistic_regression

path = get_path(__file__) + '/../..'

#D = d.trainingset_extended()
D = d.testset_extended()

cols = c('sde5', 'v11', 'e9')

a = range(int(D.shape[0]))
random.shuffle(a)

C = 1000000
num_bins = 10
bin_size = int(np.ceil(len(a)/num_bins))
bins = [a[i*bin_size:(i+1)*bin_size] for i in range(num_bins)]

X = D[:, cols]
y = D[:, c.isalert]

auc = num_bins*[0]
fpr = num_bins*[0]
tpr = num_bins*[0]
```

```

w = num_bins*[0]
b = num_bins*[0]
for i in range(num_bins):
    train = [item for j in range(num_bins) if j != i for item in bins[j]]
    test = bins[i]
    auc[i], fpr[i], tpr[i], w[i], b[i] = fit_logistic_regression(
        X[train,:], y[train,:], X[test,:], y[test,:], C=C)

def save_result(save_path):
    with open(save_path, 'w') as f:
        json.dump(get_result(), f, indent=4)

def get_result():
    wl = [a[0].tolist() for a in w]
    save_data = dict(zip(
        ['weights', 'auc', 'intercepts', 'C', 'num_bins', 'dataset_size', 'generator'],
        [wl, auc, b, C, num_bins, D.shape[0], __file__]
    ))
    return save_data

```

Code Listing A.6: https://github.com/alphabits/ford_challenge/tree/master/sessions/17-recreating-winning-entry/scripts/inferences_features_cv.py

A.12 Forward selection - Code

```

import json
import random

import numpy as np

from src.utils2 import c_ex as c, get_path, L_ex, LabelIndex
import src.dataloaders as d
from src.logistic import fit_logistic_regression

path = get_path(__file__) + '/../..'

D = d.trainingset_extended()
#Dt = d.testset_extended()

a = range(D.shape[0])
random.shuffle(a)

num_train_rows = 10000
num_test_rows = 5000

tr_rows = a[:num_train_rows]
ts_rows = a[num_train_rows:(num_train_rows+num_test_rows)]

```

```

X = D[:, 4:]
X = X[tr_rows, :]
y = D[tr_rows, c('isalert')]

Xt = D[:, 4:]
Xt = Xt[ts_rows, :]
yt = D[ts_rows, c('isalert')]

auc = np.zeros((90,90));

# Remove P3, P6, P8, V7 and V9 and
# the corresponding running features.
# See session 9 on data exploration
# for details
cc = LabelIndex(L_ex[4:])
exclude = cc('p3', 'p6', 'p8', 'v7', 'v9',
            'mp3', 'mp6', 'mp8', 'mv7', 'mv9',
            'sdp3', 'sdp6', 'sdp8', 'sdv7', 'sdv9')
candidates = [i for i in range(90) if i not in exclude]

chosen = []

for i in range(90):
    for c in candidates:
        features = chosen + [c]
        result = fit_logistic_regression(X[:, features], y, Xt[:, features], yt)
        auc[i, c] = result[0]

    chosen_feature = auc[i,:].argmax()

    if auc[i,chosen_feature] <= auc[i-1,:].max():
        break

    candidates.remove(chosen_feature)
    chosen.append(chosen_feature)

```

Code Listing A.7: https://github.com/alphabits/ford_challenge/tree/master/sessions/18-forward-selection/scripts/forward_selection.py

A.13 Forward selection - Results

Feature added	AUC	Feature added	AUC
V11	0.6978	<i>... continued</i>	
E9	0.8089	sdV3	0.8799
sdE1	0.8333	sdE9	0.8802
mE6	0.8500	sdE8	0.8811
sdE4	0.8539	mE7	0.8818
mV3	0.8566	mE9	0.8822
sdV8	0.8597	mV4	0.8827
E8	0.8612	mV5	0.8832
mE10	0.8623	mE1	0.8833
mP5	0.8635	E6	0.8835
sdP5	0.8687	V4	0.8839
mP7	0.8705	V2	0.8839
mV11	0.8717	E10	0.8840
V10	0.8727	V5	0.8841
sdV1	0.8737	sdE3	0.8843
mE4	0.8744	sdE11	0.8843
sdP2	0.8759	P4	0.8844
sdE6	0.8766	mE11	0.8846
sdP7	0.8771	sdP4	0.8848
sdP1	0.8778	sdV2	0.8848
sdE2	0.8781	sdV4	0.8849
mV2	0.8785	mV8	0.8851
mE8	0.8790	mV6	0.8855
sdV10	0.8793	mP1	0.8857
<i>continues ...</i>		sdE5	0.8858

Bibliography

- Abou-Nasr, M. (2011a). Kaggle forum: About the parameter - Reply 3. URL <http://www.kaggle.com/c/stayalert/forums/t/317/about-the-parameter/1861#post1861>.
- Abou-Nasr, M. (2011b). Kaggle forum: Discrete or continuous values - Reply 2. URL <http://www.kaggle.com/c/stayalert/forums/t/266/discrete-or-continuous-values/1635#post1635>.
- Abou-Nasr, M. (2011c). Kaggle forum: Were units changed? - Reply 2. URL <http://www.kaggle.com/c/stayalert/forums/t/268/were-units-changed/1641#post1641>.
- ACM (2011). KDD Cup Center. URL <http://www.sigkdd.org/kddcup/index.php>.
- Alpaydin, E. (2010). *Introduction to Machine Learning*. The MIT Press, 2 edition.
- Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer Science + Business Media LLC.
- David (2011). Methods/Tips from non-top 3 participants. URL <http://www.kaggle.com/c/stayalert/forums/t/328/methods-tips-from-non-top-3-participants/1967#post1967>.
- Hastie, T., Tibshirani, R., and Friedman, J. (2009). *The Elements of Statistical Learning*. Springer Science + Business Media LLC.

- Inference (2011a). First place in the stay alert! competition. URL <http://www.kaggle.com/blog/wp-content/uploads/2011/03/ford.pdf>.
- Inference (2011b). Kaggle forum: Ford - Reply 3. URL <http://www.kaggle.com/c/stayalert/forums/t/295/ford/1743#post1743>.
- Kaggle.com (2011). The Ford Challenge Data Files. URL <http://www.kaggle.com/c/stayalert/Data>.
- Labs, Y. (2011). KDD Cup 2011. URL <http://kddcup.yahoo.com/>.
- Mørup, M. (2011). Meetings with morten mørup.
- Netflix (2011). Netflix Prize Leaderboard. URL <http://www.netflixprize.com/leaderboard>.
- Pardos, Z. (2011). Kaggle forum: Top two teams with same AUC - Reply 1. URL <http://www.kaggle.com/c/stayalert/forums/t/327/top-two-teams-with-same-auc/1937#post1937>.
- Rosanne (2011). Kaggle forum: Top two teams with same AUC - Reply 4. URL <http://www.kaggle.com/c/stayalert/forums/t/327/top-two-teams-with-same-auc/1957#post1957>.
- Wasserman, L. A. (2004). *All of Statistics: a concise course in statistical inference*. Springer Science + Business Media Inc.
- Wikipedia (2011). Netflix Prize. URL http://en.wikipedia.org/wiki/Netflix_Prize.